



The plot at left represents the predictor space (on X_1 and X_2) with a training data set plotted and the class of their response variable indicated by the color.

- If we consider this a classification tree without any splits yet (i.e. only one region R_1), what would be the prediction for *every* new observation?
- What is the (training) misclassification rate?
- What is the GINI index?
- What is the cross-entropy?